

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn import preprocessing, svm
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import StandardScaler
```

```
In [2]: df=pd.read_csv(r"C:\Users\yasoda\Documents\202U1A05C1\insurance.csv")
df
```

Out[2]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

data cleaning and preprocessing

```
In [3]: df.head()
```

```
Out[3]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
In [4]: df.tail()
```

```
Out[4]:
```

	age	sex	bmi	children	smoker	region	charges
1333	50	male	30.97	3	no	northwest	10600.5483
1334	18	female	31.92	0	no	northeast	2205.9808
1335	18	female	36.85	0	no	southeast	1629.8335
1336	21	female	25.80	0	no	southwest	2007.9450
1337	61	female	29.07	0	yes	northwest	29141.3603

```
In [5]: df.shape
```

```
Out[5]: (1338, 7)
```

```
In [6]: df.describe()
```

Out[6]:


	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

```
In [7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

```
In [8]: sns.lmplot(x="age",y="charges",data=df,order=2,ci=None)
```

```
Out[8]: <seaborn.axisgrid.FacetGrid at 0x1620c656bc0>
```



In the above scatter plot graph we can able to know that the aged peoples charges are low

```
In [9]: df.fillna(method='ffill',inplace=True)
```


```
In [10]: x=np.array(df['age']).reshape(-1,1)  
y=np.array(df['charges']).reshape(-1,1)
```

```
In [11]: df.dropna(inplace=True)
```

```
In [12]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)  
regr=LinearRegression()  
regr.fit(x_train,y_train)  
print(regr.score(x_test,y_test))
```


```
0.08293882256151675
```

```
In [13]: y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='g')
plt.plot(x_test,y_pred,color='y')
plt.show()
```




```
In [14]: df500=df[:][:500]  
sns.lmplot(x="age",y="charges",data=df500,order=1,ci=None)
```

```
Out[14]: <seaborn.axisgrid.FacetGrid at 0x1623d9fe830>
```



```
In [15]: df500.fillna(method='ffill',inplace=True)
x=np.array(df500['age']).reshape(-1,1)
y=np.array(df500['charges']).reshape(-1,1)
df500.dropna(inplace=True)
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25)
regr=LinearRegression()
regr.fit(x_train,y_train)
print("Regression:",regr.score(x_test,y_test))
y_pred=regr.predict(x_test)
plt.scatter(x_test,y_test,color='r')
plt.plot(x_test,y_pred,color='g')
plt.show()
```

Regression: -0.06442340865508056




```
In [16]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
model=LinearRegression()
model.fit(x_train,y_train)
y_pred=model.predict(x_test)
r2=r2_score(y_test,y_pred)
print("R2 Score:",r2)
```

R2 Score: -0.06442340865508056

```
In [17]: df.isnull().sum()
```

```
Out[17]: age      0
sex        0
bmi        0
children   0
smoker     0
region     0
charges    0
dtype: int64
```

Implementation of Ridge regression

```
In [157]: from sklearn.linear_model import Ridge,RidgeCV,Lasso
from sklearn.preprocessing import StandardScaler
```

```
In [132]: convert={"sex":{"male":1,"female":2}}
df=df.replace(convert)
df
```

12	23	1	34.400	0	2	2	1826.843000
13	56	2	39.820	0	2	1	11090.717800
14	27	1	42.130	0	1	1	39611.757700
15	19	1	24.600	1	2	2	1837.237000
16	52	2	30.780	1	2	3	10797.336200
17	23	1	23.845	0	2	3	2395.171550
18	56	1	40.300	0	2	2	10602.385000
19	30	1	35.300	0	1	2	36837.467000
20	60	2	36.005	0	2	3	13228.846950
21	30	2	32.400	1	2	2	4149.736000
22	18	1	34.100	0	2	1	1137.011000
23	34	2	31.920	1	1	3	37701.876800
24	37	1	28.025	2	2	4	6203.901750

```
In [133]: convert={"smoker":{"yes":1,"no":2}}
df=df.replace(convert)
df
```

35	19	1	20.425	0	2	4	1625.433750
36	62	2	32.965	3	2	4	15612.193350
37	26	1	20.800	0	2	2	2302.300000
38	35	1	36.670	1	1	3	39774.276300
39	60	1	39.900	0	1	2	48173.361000
40	24	2	26.600	0	2	3	3046.062000
41	31	2	36.630	2	2	1	4949.758700
42	41	1	21.780	1	2	1	6272.477200
43	37	2	30.800	2	2	1	6313.759000
44	38	1	37.050	1	2	3	6079.671500
45	55	1	37.300	0	2	2	20630.283510
46	18	2	38.665	2	2	3	3393.356350
47	28	2	34.770	0	2	4	3556.922300

```
In [134]: convert={"region":{"southeast":3,"southwest":4,"northeast":5,"northwest":6}}
df=df.replace(convert)
df
```

Out[134]:

	age	sex	bmi	children	smoker	region	charges
0	19	2	27.900	0	1	2	16884.924000
1	18	1	33.770	1	2	1	1725.552300
2	28	1	33.000	3	2	1	4449.462000
3	33	1	22.705	0	2	4	21984.470610
4	32	1	28.880	0	2	4	3866.855200
5	31	2	25.740	0	2	1	3756.621600
6	46	2	33.440	1	2	1	8240.589600
7	37	2	27.740	3	2	4	7281.505600
8	37	1	29.830	2	2	3	6406.410700
9	60	2	25.840	0	2	4	28923.136920
10	25	1	26.220	0	2	3	2721.320800

```
In [136]: features = df.columns[0:1]
target = df.columns[-1]
#X and y values
X = df[features].values
y = df[target].values
#split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=17)
print("The dimension of x_train is {}".format(X_train.shape))
print("The dimension of x_test is {}".format(X_test.shape))
#Scale features
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

The dimension of x_train is (936, 1)
The dimension of x_test is (402, 1)

```
In [142]: ridgeReg=Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
train_score_ridge=ridgeReg.score(X_train,y_train)
test_score_ridge=ridgeReg.score(X_test,y_test)
print("\nRidge Model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.07446228994221393
The test score for ridge model is 0.10855133360950642

```
In [160]: lr = LinearRegression()
#Fit model
lr.fit(X_train, y_train)
#predict
#prediction = lr.predict(X_test)
#actual
actual = y_test
train_score_lr = lr.score(X_train, y_train)
test_score_lr = lr.score(X_test, y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 0.07447061146193878
The test score for lr model is 0.10891203216512224

```
In [ ]: plt.figure(figsize=(10,10))
add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge,$\alpha=grid$')
add plot for
```


```
In [149]: print("\nLasso Model: \n")
lasso = Lasso(alpha = 10)
lasso.fit(X_train,y_train)
train_score_ls =lasso.score(X_train,y_train)
test_score_ls =lasso.score(X_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```

Lasso Model:

The train score for ls model is 0.07446997086306062

The test score for ls model is 0.10881427793326703

```
In [156]: pd.Series(lasso.coef_, features).sort_values(ascending = True).plot(kind = "bar")  
plt.show()
```




```
In [152]: from sklearn.linear_model import LassoCV  
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(X_train,y_train)  
print(lasso_cv.score(X_train,y_train))  
print(lasso_cv.score(X_test,y_test))
```

0.07446997086306062

0.10881427793326703

```
In [165]: plt.figure(figsize=(10,10))
#add plot for ridge regression
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label='Ridge;$\alpha=10$',border=1)
#add plot for lasso regression
plt.plot(lasso_cv.coef_,alpha=0.6,linestyle='none',marker='d',markersize=6,color='blue',label=r'Ridge;$\alpha=grid$')
#add plot for linear model
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='Linear Regression')
#rotate axis
plt.xticks(rotation=90)
plt.legend()
plt.title("comparison plot of Ridge, Lasso and Linear regression model")
plt.show()
```




```
In [169]: from sklearn.linear_model import RidgeCV
#Ridge Cross validation
ridge_cv = RidgeCV(alphas = [0.0001, 0.001,0.01, 0.1, 1, 10]).fit(X_train, y_train)
#score
print(ridge_cv.score(X_train,y_train))
print(ridge_cv.score(X_test,y_test))
```

```
0.07446228994221393
0.10855133360950775
```

Elastic net regression

```
In [51]: from sklearn.linear_model import ElasticNet
regr=ElasticNet()
regr.fit(x,y)
print(regr.coef_)
print(regr.intercept_)
```

```
[283.49314745]
[2117.95065707]
```

```
In [52]: y_pred_elastic=regr.predict(x_train)
```

```
In [53]: mean_squared_error=np.mean((y_pred_elastic-y_train)**2)
print("Mean Squared Error on test set",mean_squared_error)
```

```
Mean Squared Error on test set 161136962.85931197
```

Logistic regression

```
In [54]: import pandas as pd
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.preprocessing import StandardScaler
```

```
In [55]: df=pd.read_csv(r"C:\Users\yasoda\Documents\202U1A05C1\insurance.csv")
df
```

Out[55]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [56]: pd.set_option('display.max_rows',10000000000)
pd.set_option('display.max_columns',10000000000)
pd.set_option('display.width',95)
```

```
In [57]: print('This DataFrame has %d Rows and %d columns'%(df.shape))
```

This DataFrame has 1338 Rows and 7 columns

```
In [58]: convert={"smoker":{"yes":1,"no":2}}
df=df.replace(convert)
df
```

Out[58]:

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	1	southwest	16884.924000
1	18	male	33.770	1	2	southeast	1725.552300
2	28	male	33.000	3	2	southeast	4449.462000
3	33	male	22.705	0	2	northwest	21984.470610
4	32	male	28.880	0	2	northwest	3866.855200
5	31	female	25.740	0	2	southeast	3756.621600
6	46	female	33.440	1	2	southeast	8240.589600
7	37	female	27.740	3	2	northwest	7281.505600
8	37	male	29.830	2	2	northeast	6406.410700
9	60	female	25.840	0	2	northwest	28923.136920
10	25	male	26.220	0	2	northeast	2721.320800

```
In [59]: convert={"sex":{"male":8,"female":9}}
df=df.replace(convert)
df
```

1305	24	9	27.720	0	2	southeast	2464.618800
1306	29	9	21.850	0	1	northeast	16115.304500
1307	32	8	28.120	4	1	northwest	21472.478800
1308	25	9	30.200	0	1	southwest	33900.653000
1309	41	8	32.200	2	2	southwest	6875.961000
1310	42	8	26.315	1	2	northwest	6940.909850
1311	33	9	26.695	0	2	northwest	4571.413050
1312	34	8	42.900	1	2	southwest	4536.259000
1313	19	9	34.700	2	1	southwest	36397.576000
1314	30	9	23.655	3	1	northwest	18765.875450
1315	18	8	28.310	1	2	northeast	11272.331390
1316	19	9	20.600	0	2	southwest	1731.677000
1317	18	8	53.130	0	2	southeast	1163.462700

```
In [60]: features_matrix=df.iloc[:,0:4]
```

```
In [61]: target_vector=df.iloc[:, -3]
```

```
In [170]: print('The Features Matrix Has %d Rows And %d Column(s)'%(features_matrix.shape))
```

The Features Matrix Has 1338 Rows And 4 Column(s)

```
In [171]: print('The Target Matrix Has %d Rows And %d Column(s)'%(np.array(target_vector).reshape(-1,1).shape))
```

The Target Matrix Has 1338 Rows And 1 Column(s)

```
In [64]: features_matrix_standardized=StandardScaler().fit_transform(features_matrix)
```

```
In [177]: scaling=1,class_weight=None,random_state=None,solver='lbfgs',max_iter=100,multi_class='auto',verbose=0,warm_start=False,n_jobs=
```

```
In [180]: Logistic_Regression_Model=algorithm.fit(features_matrix_standardized,target_vector)
```

```
In [181]: observation=[[1,0,0.99539,-0.05889,]]
```

```
In [183]: predictions=Logistic_Regression_Model.predict(observation)
print('The model Predicted the observation to belong to class %s'%(predictions))
```

The model Predicted the observation to belong to class [2]

```
In [184]: print('The algorithm was trained to predict one of the two classes: %s'%(algorithm.classes_))
```

The algorithm was trained to predict one of the two classes: [1 2]

```
In [185]: print(" " "The model says she probability of the observation we passed belonging to class[0] Is %s" " "%(algorithm.predict_proba(observation)[0][0]))
```

The model says she probability of the observation we passed belonging to class[0] Is 0.1942921563693959

```
In [71]: print()
```

```
In [186]: says the probabaility of the observation we passed belonging to class[1] Is %s" " "%(algorithm.predict_proba(observation)[0][1]))
```

The model says the probabaility of the observation we passed belonging to class[1] Is 0.1942921563693959

```
In [73]: x=np.array(df['age']).reshape(-1,1)
y=np.array(df['smoker']).reshape(-1,1)
```

```
In [74]: lerg=LogisticRegression()
lerg.fit(x,y)
print(lerg.score(x,y))
```

0.7952167414050823

C:\Users\yasoda\AppData\Local\Programs\Python\Python310\lib\site-packages\sklearn\utils\validation.py:1143: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

Decision tree regression

```
In [75]: import numpy as np
import pandas as pd
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
```

```
In [76]: df=pd.read_csv(r"C:\Users\yasoda\Documents\202U1A05C1\insurance.csv")
df
```

1080	18	male	21.780	2	no	southeast	11884.048580
1081	32	male	27.835	1	no	northwest	4454.402650
1082	38	male	19.950	1	no	northwest	5855.902500
1083	32	male	31.500	1	no	southwest	4076.497000
1084	62	female	30.495	2	no	northwest	15019.760050
1085	39	female	18.300	5	yes	southwest	19023.260000
1086	55	male	28.975	0	no	northeast	10796.350250
1087	57	male	31.540	0	no	northwest	11353.227600
1088	52	male	47.740	1	no	southeast	9748.910600
1089	56	male	22.100	0	no	southwest	10577.087000
1090	47	male	36.190	0	yes	southeast	41676.081100
1091	55	female	29.830	0	no	northeast	11286.538700
1092	23	male	32.700	3	no	southwest	3591.480000

```
In [77]: df['region'].value_counts()
```

```
Out[77]: region
southeast    364
southwest    325
northwest    325
northeast    324
Name: count, dtype: int64
```

```
In [78]: df['bmi'].value_counts()
```

```
28.580    1  
24.090    1  
25.100    1  
34.300    1  
43.400    1  
39.200    1  
35.700    1  
26.070    1  
39.425    1  
40.480    1  
38.900    1  
47.410    1  
35.435    1  
46.700    1  
46.200    1  
23.800    1  
44.770    1  
32.120    1  
30.970    1  
Name: count, dtype: int64
```



```
In [79]: convert={"sex":{"male":1,"female":0}}
df=df.replace(convert)
df
```

3	33	1	22.705	0	no	northwest	21984.470610
4	32	1	28.880	0	no	northwest	3866.855200
5	31	0	25.740	0	no	southeast	3756.621600
6	46	0	33.440	1	no	southeast	8240.589600
7	37	0	27.740	3	no	northwest	7281.505600
8	37	1	29.830	2	no	northeast	6406.410700
9	60	0	25.840	0	no	northwest	28923.136920
10	25	1	26.220	0	no	northeast	2721.320800
11	62	0	26.290	0	yes	southeast	27808.725100
12	23	1	34.400	0	no	southwest	1826.843000
13	56	0	39.820	0	no	southeast	11090.717800
14	27	1	42.130	0	yes	southeast	39611.757700
15	19	1	24.600	1	no	southwest	1837.237000

```
In [80]: x=["bmi","children"]
y=["yes","no"]
all_inputs=df[x]
all_classes=df["sex"]
```

```
In [187]: (x_train,x_test,y_train,y_test)=train_test_split(all_inputs,all_classes,test_size=0.03)
```

```
In [188]: clf=DecisionTreeClassifier(random_state=0)
```

```
In [189]: clf.fit(x_train,y_train)
```

```
Out[189]:
```

▼

DecisionTreeClassifier

DecisionTreeClassifier(random_state=0)

```
In [190]: score=clf.score(x_test,y_test)
print(score)
```

0.4878048780487805

Random forest

```
In [107]: import pandas as pd
import numpy as ny
import matplotlib.pyplot as plt,seaborn as sns
```

```
In [108]: df=pd.read_csv(r"C:\Users\yasoda\Documents\202U1A05C1\insurance.csv")
df
```

446	60	male	29.640	0	no	northeast	12730.999600
447	56	female	25.650	0	no	northwest	11454.021500
448	40	female	29.600	0	no	southwest	5910.944000
449	35	male	38.600	1	no	southwest	4762.329000
450	39	male	29.600	4	no	southwest	7512.267000
451	30	male	24.130	1	no	northwest	4032.240700
452	24	male	23.400	0	no	southwest	1969.614000
453	20	male	29.735	0	no	northwest	1769.531650
454	32	male	46.530	2	no	southeast	4686.388700
455	59	male	37.400	0	no	southwest	21797.000400
456	55	female	30.140	2	no	southeast	11881.969600
457	57	female	30.495	0	no	northwest	11840.775050
458	56	male	39.600	0	no	southwest	10601.412000

```
In [109]: df['charges'].value_counts()
```

```
7758.288200    1
7512.267000    1
11840.775050    1
11881.969600    1
21797.000400    1
4686.388700    1
1769.531650    1
1969.614000    1
4032.240700    1
4762.329000    1
37079.372000    1
5910.944000    1
11454.021500    1
12730.999600    1
7345.084000    1
26109.329050    1
28287.897660    1
1149.395900    1
29141.360300    1
Name: count, dtype: int64
```

```
In [110]: m={"region":{"southeast":1,"southwest":2,"northeast":3,"northwest":4}}
df=df.replace(m)
print(df)
```

370	61	female	21.090	0	no	4	13415.038100
371	57	female	22.230	0	no	3	12029.286700
372	42	female	33.155	1	no	3	7639.417450
373	26	male	32.900	2	yes	2	36085.219000
374	20	male	33.330	0	no	1	1391.528700
375	23	female	28.310	0	yes	4	18033.967900
376	39	female	24.890	3	yes	3	21659.930100
377	24	male	40.150	0	yes	1	38126.246500
378	64	female	30.115	3	no	4	16455.707850
379	62	male	31.460	1	no	1	27000.984730
380	27	female	17.955	2	yes	3	15006.579450
381	55	male	30.685	0	yes	3	42303.692150
382	55	male	33.000	0	no	1	20781.488920
383	35	female	43.340	2	no	1	5846.917600
384	44	male	22.135	2	no	3	8302.535650
385	19	male	34.400	0	no	2	1261.859000
386	58	female	39.050	0	no	1	11856.411500
387	50	male	25.365	2	no	4	30284.642940
388	26	female	22.610	0	no	4	3176.815900
389	24	female	30.210	3	no	4	4618.079900

```
In [111]: df.shape
```

```
Out[111]: (1338, 7)
```

```
In [191]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[191]: ▾ RandomForestClassifier
RandomForestClassifier()
```

```
In [194]: rf=RandomForestClassifier()
```

```
In [196]: params={'max_depth':[2,3,5,10,20], 'min_samples_leaf':[5,10,20,50,100,200], 'n_estimators':[10,25,30,50,100,200]}
```

```
In [ ]: from sklearn.model_selection import GridSearchCV  
grid_search=GridSearchCV(estimator=rf,param_grid=params,cv=2,scoring="accuracy")  
grid_search.fit(x_train,y_train)
```

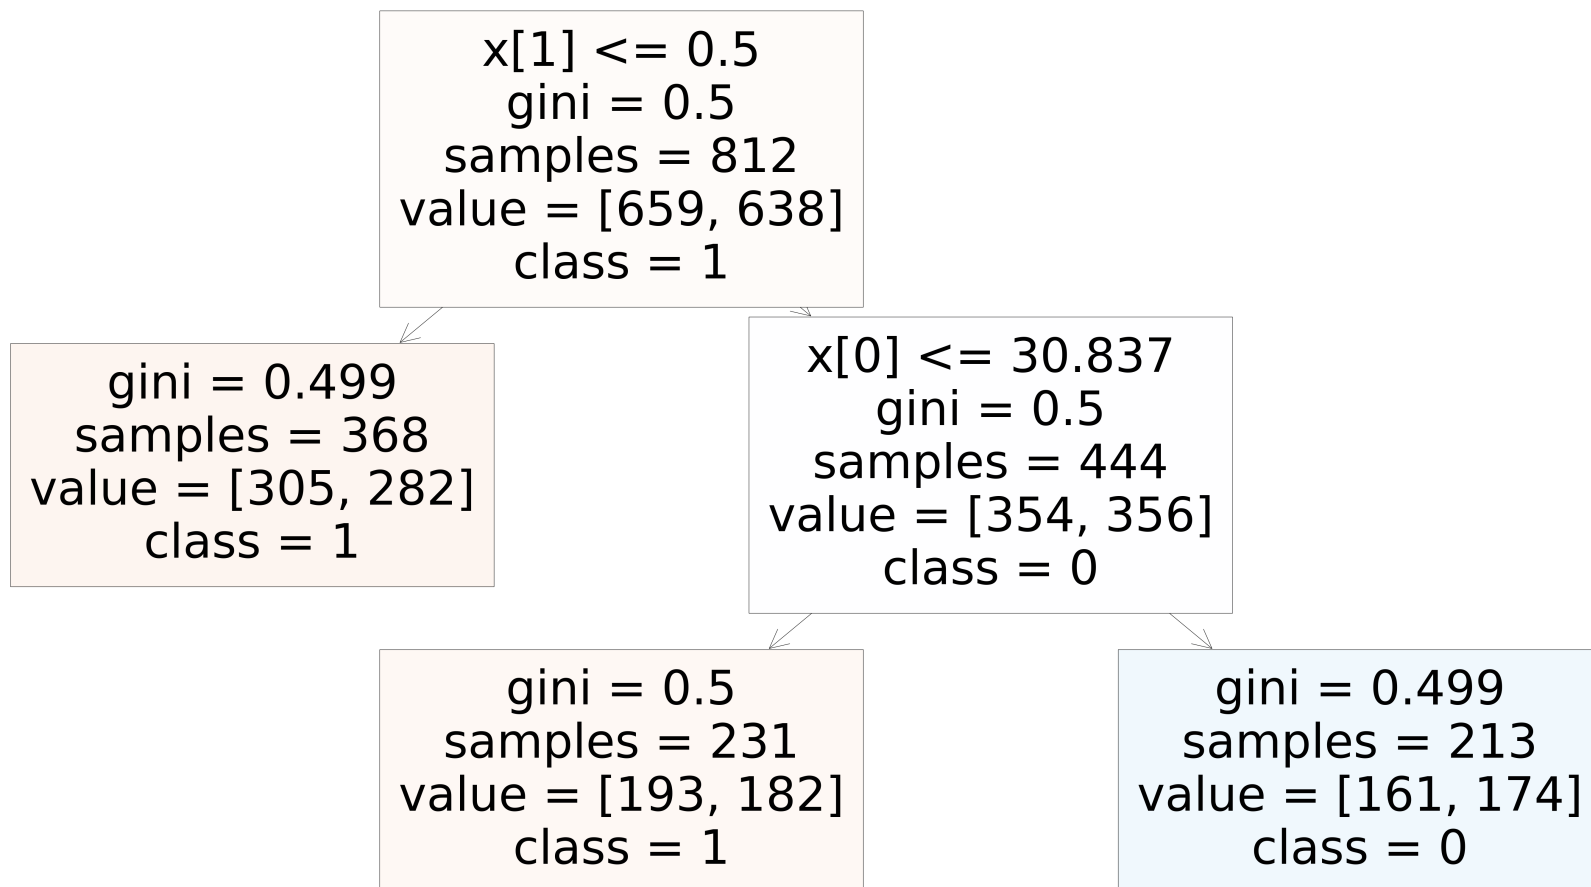
```
In [203]: grid_search.best_score_
```

```
Out[203]: 0.5250444653693241
```

```
In [204]: rf_best=grid_search.best_estimator_  
print(rf_best)
```

```
RandomForestClassifier(max_depth=20, min_samples_leaf=200, n_estimators=25)
```

```
In [205]: from sklearn.tree import plot_tree
plt.figure(figsize=(80,40))
plot_tree(rf_best.estimators_[4],class_names=['1','0'],filled=True);
```



```
In [210]: rf_best.feature_importances_
```

```
Out[210]: array([0.64289857, 0.35710143])
```

```
In [209]: score=rfc.score(x_test,y_test)
          print(score)
```

0.43902439024390244

Conclusion:In a given dataset we have performed Linear,Logistic,Decisiontree,Randomforest regression models have concluded the most accuracy it is occurred in logistic regression.when we compare other regression models Logistic Regression suits best for the given data